

# **Learning Group Structured Representations**


CS7180 Geometric DL

John Park



March 23 2023

# Bio

Reinforcement learning and symmetry







Strong inductive biases to speed up learning and improve generalization



# Symmetry


Important geometric prior



## References

1. <https://developer.nvidia.com/blog/deploying-a-scalable-object-detection-pipeline-the-inferencing-process-part-2/>
2. <https://robdhess.github.io/Steerable-E3-GNN/>
3. Monti, Federico, et al. "Geometric deep learning on graphs and manifolds using mixture model cnns." CVPR 2017.
4. Bronstein, Michael M., et al. "Geometric deep learning: going beyond euclidean data." IEEE Signal Processing Magazine 34.4 (2017): 18-42.
5. <https://github.com/sukjulian/coronary-mesh-convolution>


# Equivariance




$$f(\rho_{in}(g)x) = \rho_{out}(g)f(x)$$

# Equivariance

CNN



E(2)-CNN



## References

1. <https://github.com/QUVA-Lab/e2cnn>

# Limitations

$$f(\rho_{in}(g)x) = \rho_{out}(g)f(x)$$

Symmetry group needs to be known

# Limitations

$$f(\rho_{in}(g)x) = \rho_{out}(g)f(x)$$

Symmetry group needs to be known

Group action on input and output spaces must be known

# Limitations


$$f(\rho_{in}(g)x) = \rho_{out}(g)f(x)$$

Symmetry group needs to be known

Group action on input and output spaces must be known


Equivariant model tied to group and action

# Transformations in latent space




4 actions: 



# Transformations in latent space




4 actions:



# Transformations in latent space



4 actions: 



The mug is **on the right of** the robot gripper



The mug is **in front of** the robot gripper



# Learning symmetry from data

**Learning the Irreducible Representations of Commutative Lie Groups**

---

Taco Cohen  
Max Welling  
Machine Learning Group, University of Amsterdam

**Abstract**

We present a new probabilistic model of compact commutative Lie groups that produces invariant-equivariant and disentangled representations of data. To define the notion of disentangling, we borrow a fundamental principle from physics that is used to derive the elementary particles of a system from its symmetries. Our model employs a newfound Bayesian conjugacy relation that enables fully tractable probabilistic inference over compact commutative Lie groups – a class that includes the groups that describe the rotation and cyclic translation of images. We train the model on pairs of transformed image patches, and show that the learned invariant representation is highly effective for classification.

**1. Introduction**

Recently, the field of deep learning has produced some remarkable breakthroughs. The hallmark of the deep learning approach is to learn multiple layers of *representation* of data, and much work has gone into the develop-

T.S.COHEN@UVA.NL  
M.WELLING@UVA.NL

for meaning for future work.

What do we mean, intuitively, when we speak of invariance and disentanglement? A disentangled representation is one that explicitly represents the distinct factors of variation in the data. For example, visual data (i.e. pixels) can be thought of as a composition of object identity, position and pose, lighting conditions, etc. Once disentangling is achieved, invariance follows easily: to build a representation that is invariant to the transformation of a factor of variation (e.g. object position) that is considered a nuisance for a particular task (e.g. object classification), one can simply ignore the units in the representation that encode the nuisance factor.

To get a mathematical handle on the concept of disentangling, we borrow a fundamental principle from physics, which we refer to as Weyl's principle, following Kanatani (1990). In physics, this idea is used to tease apart (i.e. disentangle) the *elementary particles* of a physical system from mere measurement values that have no inherent physical significance. We apply this principle to the area of vision, after all, pixels are nothing but physical measurements.

Weyl's principle measures over a commutative group that acts

# Learning symmetry from data

**Learning the Irreducible Lie**

Taco Cohen  
Max Welling  
Machine Learning Group, University of Amsterdam

**Abstract**

We present a new probabilistic model of commutative Lie groups that produces invariant, equivariant and disentangled representations of data. To define the notion of disentangling, we borrow a fundamental principle from physics: the elementary particles of a system are defined by their symmetries. Our model employs a newfound Bayesian conjugacy relation that enables fully tractable probabilistic inference over compact commutative Lie groups – a class that includes the groups that describe the rotations and cyclic translations of images. We train the model on pairs of transformed image patches, and show that the learned invariant representation is effective for classification.

**1. Introduction**

Recently, the field of deep learning has produced remarkable breakthroughs. The hallmark of the learning approach is to learn multiple layers of *invariance* of data, and much work has gone into finding

**Learning Lie Groups for Invariant Visual Perception\***

Rajesh P. N. Rao and Daniel L. Ruderman  
Sloan Center for Theoretical Neurobiology  
The Salk Institute  
La Jolla, CA 92037  
[{rao, ruderman}@salk.edu](mailto:{rao, ruderman}@salk.edu)

**Abstract**

One of the most important problems in visual perception is that of visual invariance: how are objects perceived to be the same despite undergoing transformations such as translations, rotations or scaling? In this paper, we describe a Bayesian method for learning invariances based on Lie group theory. We show that our approach is based on first principles. The key idea is that certain invariance classes can be regarded as special cases of the Lie group approach, the latter being capable of handling in principle arbitrarily large transformations. Using a matrix-exponential based generative model of images, we derive an unsupervised algorithm for learning Lie group operators from input data containing infinitesimal transformations. The on-line unsupervised learning algorithm maximizes the posterior probability of generating the training data. We provide experimental results suggesting that the proposed method can learn Lie group operators for handling reasonably large 1-D translations and 2-D rotations.

**1 INTRODUCTION**

A fundamental problem faced by both biological and machine vision systems is the recognition of familiar objects and patterns in the presence of transformations such as translations, rotations and scaling. The importance of this problem was recognized early by visual scientists such as J. J. Gibson who hypothesized that “constant perception depends on the ability of the individual to detect the invariants” [6]. Among computational neuroscientists, Pitts and McCulloch were perhaps the first to propose a method for perceptual invariance (“knowing universals”) [12]. A number of other approaches have since been proposed [5, 7, 10], some relying on temporal sequences of input patterns undergoing transformations (e.g. [4]) and others relying on modifications to the distance

# Learning symmetry from data

**Learning the Irreducible Lie Algebra**

Taco Cohen  
Max Welling  
Machine Learning Group, University of Amsterdam

**Abstract**

We present a new probabilistic model of commutative Lie groups that produces invariant equivariant and disentangled representations of data. To define the notion of disentangling we borrow a fundamental principle from physics: that particles do not interact with each other. This is used to derive the elementary particles of a system from its symmetries. Our model employs a newfound Bayesian conjugacy relation that makes fully tractable probabilistic inference over compact commutative Lie groups – a class that includes the groups that describe the rotations and cyclic translations of images. We train the model on pairs of transformed image patches, and show that the learned invariant representation is effective for classification.

**1. Introduction**

Recently, the field of deep learning has produced remarkable breakthroughs. The hallmark of the learning approach is to learn multiple layers of *invariance* of data, and much work has gone into finding

**Learning Lie Groups**

Rajesh P. N. Rao  
Sloan Center for Computational Neuroscience  
University of Southern California  
{rao, rao}@usc.edu

**Abstract**

One of the most important problems in computer vision is how to learn invariant representations such as translations, rotations, and scaling. The Bayesian method for learning invariant representations proposed here can be regarded as a special case capable of handling in principle a wide range of transformations. The one-dimensional exponential based generative algorithm for learning Lie group invariant representations. The one-dimensional posterior probability of general results suggesting that the proposed method handles reasonably large 1-D

**Automatic Symmetry Discovery with Lie Algebra Convolutional Network**

Nima Dehghani  
Northwestern University  
nimehdi@bu.edu

Robin Walters  
Northeastern University  
rwalters@northeastern.edu

Yanchen Liu  
Northeastern University  
liu.yanc@northeastern.edu


Dashun Wang  
Northwestern University  
dashun.wang@kellogg.northwestern.edu

Rose Yu  
University of California San Diego  
roseyu@ucsd.edu


**Abstract**

Existing equivariant neural networks require prior knowledge of the symmetry group and discretization for learning invariance. We propose to work with Lie algebras (infinitesimal generators) instead of Lie groups. Our model, the Lie algebra convolutional network (L-conv) can automatically discover symmetries and does not require discretization of the group. We show that L-conv can serve as a building block to construct any group equivariant feedforward architecture. Both CNNs and Graph Convolutional Networks can be expressed as L-conv with appropriate


# Learning symmetry from data



# Learning symmetry from data



# Learning symmetry from data



# Outline

Group-structured Representations (Higgins et al. 2018)

Learning disentangled representations (Quessard et al. 2020)

Learning symmetric representations (Park et al. 2022)

# Representations

What is a good representation?

- Makes learning task easier
- Preserve underlying structure/information of data

## References

1. Bengio, Y., Courville, A., & Vincent, P. (2013). Representation learning: A review and new perspectives. *IEEE transactions on pattern analysis and machine intelligence*, 35(8), 1798-1828.
2. Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep learning. MIT press.
3. Locatello, F., Bauer, S., Lucic, M., Raetsch, G., Gelly, S., Schölkopf, B., & Bachem, O. (2019, May). Challenging common assumptions in the unsupervised learning of disentangled representations. *ICML*.
4. Higgins, I., Amos, D., Pfau, D., Racaniere, S., Matthey, L., Rezende, D., & Lerchner, A. (2018). Towards a definition of disentangled representations. *arXiv preprint arXiv:1812.02230*.

# Representations

What is a good representation?

- Makes learning task easier
- Preserve underlying structure/information of data

Properties

- Lower-dimensional → faster learning
- Useful for other downstream tasks
- Disentangled (separate factors of variation)

References


1. Bengio, Y., Courville, A., & Vincent, P. (2013). Representation learning: A review and new perspectives. *IEEE transactions on pattern analysis and machine intelligence*, 35(8), 1798-1828.
2. Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep learning. MIT press.
3. Locatello, F., Bauer, S., Lucic, M., Raetsch, G., Gelly, S., Schölkopf, B., & Bachem, O. (2019, May). Challenging common assumptions in the unsupervised learning of disentangled representations. *ICML*.
4. Higgins, I., Amos, D., Pfau, D., Racaniere, S., Matthey, L., Rezende, D., & Lerchner, A. (2018). Towards a definition of disentangled representations. *arXiv preprint arXiv:1812.02230*.

# Disentangled Representations

Good representations should disentangle  
the **factors of the data generating  
process**

- Color
- Shape
- Background

No agreement on exact definition



## References

1. Van Steenkiste, S., Locatello, F., Schmidhuber, J., & Bachem, O. (2019). Are disentangled representations helpful for abstract visual reasoning?. NeurIPS

# Disentangled representations

## Towards a Definition of Disentangled Representations

Irina Higgins\*, David Amos\*, David Pfau, Sebastien Racaniere,  
Loic Matthey, Danilo Rezende, Alexander Lerchner  
DeepMind

{irinah,davidamos,pfau,sracaniere,  
lmatthey,danilor,lerchner}@google.com

December 7, 2018

How can intelligent agents solve a diverse set of tasks in a data-efficient manner? The disentangled representation learning approach posits that such an agent would benefit from separating out (disentangling) the underlying structure of the world into disjoint parts of its representation. However, there is no generally agreed-upon definition of disentangling, not least because it is unclear how to formalise the notion of world structure beyond toy datasets with a known ground truth generative process. Here we propose that a principled solution to characterising disentangled representations can be found


Connects disentanglement with symmetry  
inspired by physics

Shift focus from objects to transformations

Observe how an object transforms to learn  
disentangled representation

# Disentangled Representation

A vector representation is called a **disentangled representation** with respect to a particular decomposition of a symmetry group into subgroups, if it decomposes into independent subspaces, where each subspace is affected by the action of a single subgroup, and the actions of all other subgroups leave the subspace unaffected.



$$G = G_x \times G_y \times G_c$$

$$Z = Z_x \times Z_y \times Z_c$$

# Paper 1

---

## Learning Disentangled Representations and Group Structure of Dynamical Environments

---

Robin Quessard<sup>1,2\*</sup> Thomas D. Barrett<sup>3\*</sup> William R. Clements<sup>1\*</sup>

<sup>1</sup>Indust.ai, Paris, France

<sup>2</sup>École Normale Supérieure, Paris, France

<sup>3</sup>University of Oxford, Oxford, UK

robin.quessard@ens.fr, thomas.barrett@physics.ox.ac.uk, william.clements@indust.ai

### Abstract

Learning disentangled representations is a key step towards effectively discovering and modelling the underlying structure of environments. In the natural sciences, physics has found great success by describing the universe in terms of symmetry preserving transformations. Inspired by this formalism, we propose a framework, built upon the theory of group representation, for learning representations of a dynamical environment structured around the transformations that generate its evolution. Experimentally, we learn the structure of explicitly symmetric environments without supervision from observational data generated by sequential interactions. We further introduce an intuitive disentanglement regularisation to ensure the interpretability of the learnt representations. We show that our method enables accurate long-horizon predictions, and demonstrate a correlation between the quality of predictions and disentanglement in the latent space.

Practical method to learn representations using definition from Higgins et al. 2018.

Focuses on  $SO(n)$

Considers a trajectory of objects and transformations

# Method

Dataset:  $(o_1, a_1, o_2 \dots, a_T, o_T)$

Goal: Learn encoder  $f_\phi : O \rightarrow V$ ,  $a_i \mapsto g_i \in GL(V)$ , decoder  $d_\psi : V \mapsto O$

$$\hat{o}_{i+m}(\phi, \psi, \theta) = d_\psi(g_{i+m}(\theta).g_{i+m-1}(\theta) \dots g_{i+1}(\theta).f_\phi(o_i))$$

# Method

Parameterize  $g$  as product of rotations

$$g(\theta_{1,2}, \theta_{1,3}.., \theta_{n-1,n}) = \prod_{i=1}^{n-1} \prod_{j=i+1}^n R_{i,j}(\theta_{i,j})$$

$$R = R_z(\alpha) R_y(\beta) R_x(\gamma) = \begin{bmatrix} \cos \alpha & -\sin \alpha & 0 \\ \sin \alpha & \cos \alpha & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \text{yaw} & & \\ & \cos \beta & 0 & \sin \beta \\ & 0 & 1 & 0 \\ & -\sin \beta & 0 & \cos \beta \end{bmatrix} \begin{bmatrix} \text{pitch} & & \\ & 1 & 0 & 0 \\ & 0 & \cos \gamma & -\sin \gamma \\ & 0 & \sin \gamma & \cos \gamma \end{bmatrix}$$

# Method

$$\hat{o}_{i+m}(\phi, \psi, \theta) = d_\psi(g_{i+m}(\theta) \cdot g_{i+m-1}(\theta) \dots g_{i+1}(\theta) \cdot f_\phi(o_i))$$

Learn parameters  $\phi, \psi, \theta$


$$\mathcal{L}_{total} = \mathcal{L}_{rec} + \lambda \mathcal{L}_{ent}$$

$$\mathcal{L}_{rec} = l(o_{i+1}, \hat{o}_{i+1})$$

$$\mathcal{L}_{ent}(\theta) = \sum_a \sum_{(i,j) \neq (\alpha,\beta)} |\theta_{i,j}^a|^2 \quad \text{with} \quad \theta_{\alpha,\beta}^a = \max_{i,j}(|\theta_{i,j}^a|)$$







Minimize entanglement (promotes sparsity)

# Experiments (Flatland)



# Experiments (Flatland)

$$G = C_p \times C_p$$



# Paper 2

## Learning Symmetric Embeddings for Equivariant World Models

Jung Yeon Park <sup>\*1</sup> Ondrej Biza <sup>\*1</sup> Linfeng Zhao <sup>1</sup> Jan Willem van de Meent <sup>2,1</sup> Robin Walters <sup>1</sup>

### Abstract

Incorporating symmetries can lead to highly data-efficient and generalizable models by defining equivalence classes of data samples related by transformations. However, characterizing how transformations act on input data is often difficult, limiting the applicability of equivariant models. We propose learning symmetric embedding networks (SENs) that encode an input space (e.g. images), where we do not know the effect of transformations (e.g. rotations), to a feature space that transforms in a known manner under these operations. This network can be trained end-to-end with an equivariant task network to learn an explicitly symmetric representation. We validate this approach in the context of equivariant transition models with 3 distinct forms of symmetry. Our experiments demonstrate that SENs facilitate the application of equivariant networks to data with complex symmetry representations. Moreover, doing so can yield improvements in accuracy and generalization relative to both fully-equivariant and non-equivariant baselines.




Figure 1: Equivariant networks consider transformations of inputs that are easy to compute, such as in-plane rotations for MNIST digits. This paper considers transformations that are difficult to compute, such as rotations of 3D objects like cars (Carvana, 2017). Symmetric embedding networks learn representations which are transformed simply.

Kondor & Trivedi, 2018; Bao & Song, 2019; Worrall et al., 2017). This results in models that are often more parameter

Ignores disentanglement

Arbitrary groups and domains

Actions do not need to be tied to  
transformations

# Paper 2

## Learning Symmetric Embeddings for Equivariant World Models

Jung Yeon Park <sup>\*1</sup> Ondrej Biza <sup>\*1</sup> Linfeng Zhao <sup>1</sup> Jan Willem van de Meent <sup>2,1</sup> Robin Walters <sup>1</sup>

### Abstract

Incorporating symmetries can lead to highly data-efficient and generalizable models by defining equivalence classes of data samples related by transformations. However, characterizing how transformations act on input data is often difficult, limiting the applicability of equivariant models. We propose learning symmetric embedding networks (SENs) that encode an input space (e.g. images), where we do not know the effect of transformations (e.g. rotations), to a feature space that transforms in a known manner under these operations. This network can be trained end-to-end with an equivariant task network to learn an explicitly symmetric representation. We validate this approach in the context of equivariant transition models with 3 distinct forms of symmetry. Our experiments demonstrate that SENs facilitate the application of equivariant networks to data with complex symmetry representations. Moreover, doing so can yield improvements in accuracy and generalization relative to both fully-equivariant and non-equivariant baselines.




Figure 1: Equivariant networks consider transformations of inputs that are easy to compute, such as in-plane rotations for MNIST digits. This paper considers transformations that are difficult to compute, such as rotations of 3D objects like cars (Carvana, 2017). Symmetric embedding networks learn representations which are transformed simply.

Kondor & Trivedi, 2018; Bao & Song, 2019; Worrall et al., 2017). This results in models that are often more parameter


Ignores disentanglement

Arbitrary groups and domains

Actions do not need to be tied to  
transformations


# Symmetric Embedding Network (SEN)

Difficult to characterize how symmetry acts on input data, i.e. group action is unknown




# Symmetric Embedding Network (SEN)

**Key idea: Pair SEN with downstream equivariant networks and train end-to-end**




# Meta-Architecture

Learn world models with symmetry



# Meta-Architecture


Learn world models with symmetry



feature: known group action, network: equivariant neural network

# Meta-Architecture


Learn world models with symmetry



feature: known group action, network: equivariant neural network



# Meta-Architecture

Learn world models with symmetry



feature: known group action, network: equivariant neural network

# Domains



2D Shapes & 3D Blocks

$\pi/2$  rotation  
Object permutation

$C_4$ - equivariant MPNN<sup>1,2</sup>

Rush Hour

$\pi/2$  rotation  
Object permutation

$C_4$ - equivariant MPNN<sup>1,2</sup>

Reacher

$\pi/2$  rotation  
flip, translation

E(2)-steerable CNN,  
equivariant MLP<sup>3</sup>

3D Teapot

3D rotation

Matrix Multiplication

1. Cohen, Taco, and Max Welling. "Group equivariant convolutional networks." *International conference on machine learning* (2016).

2. Scarselli, Franco, et al. "The graph neural network model." *IEEE transactions on neural networks* 20.1 (2008): 61-80.

3. Weiler, Maurice, and Gabriele Cesa. "General e (2)-equivariant steerable cnns." *Advances in Neural Information Processing Systems* 32 (2019).


# Model Performance

3D Teapot

Ground truth



Non-equivariant



Ours (SEN)




# Model Performance

3D Teapot

Ground truth





Non-equivariant



Ours (SEN)



Reacher



# Generalization over limited actions

Train only on {up,right,down}, test on all {up,right,down,left}

	Limited Actions	Model	MRR (10 step, %)	DIE (10 step)
3D Blocks	{up,right,down}	Non-equivariant	$61.8 \pm 13.0$	$181 \pm 79.0$
		Fully-equivariant	$86.0 \pm 31.0$	$15 \pm 9.1$
		Ours	<b><math>100 \pm 0.0</math></b>	<b><math>5 \pm 4.7</math></b>

SEN achieves higher performance, lower equivariance error than baselines

Generalization benefits of equivariant neural networks

# Summary

## Group-structured representations

- Learn in latent space
- Connection with disentanglement

Paper 1: Learn autoencoder and parameterize transformation, use entanglement loss

Paper 2: Pair conventional network with equivariant network and train end-to-end